

Advanced Data Analysis

DATA 71200

Class 9: Naive Bayes and Decision Trees

Naive Bayes

- ▶ **Uses Bayes rule to predict the probability of an example belonging to a given class**
- ▶ **“Naive” because it examines each feature individually**
- ▶ **Calculate per-class statistics for each feature**
 - Bernoulli (binary data) - number of times a feature is non-zero for each class
 - Multinomial (count data) - average of how many times a feature occurs for each class
 - Gaussian (continuous data) - mean and standard deviation of the value of the feature for each class

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$

A, B = events

$P(A|B)$ = probability of A given B is true

$P(B|A)$ = probability of B given A is true

$P(A), P(B)$ = the independent probabilities of A and B

Naive Bayes

▸ **Assumptions**

- Independence between features
- Gaussian version assumes normally distributed data with the same variance in each class

▸ **Best Practices**

- Bernoulli and Multinomial best used on sparse data
- Gaussian version best used on high-dimensional data

Naive Bayes

▸ Parameters

- Alpha (Bernoulli and Multinomial) - model complexity
 - The larger the alpha, the more virtual data points that are added for smoothing

▸ Strengths

- Faster to train and predict than linear models
- Understandable training procedure
- Works well on large datasets
 - More efficient on them than linear models

▸ Weaknesses

- Coefficients not easily interpreted
- Does not generalize as well as linear models

Decision Trees

- Learned hierarchy of if/else questions

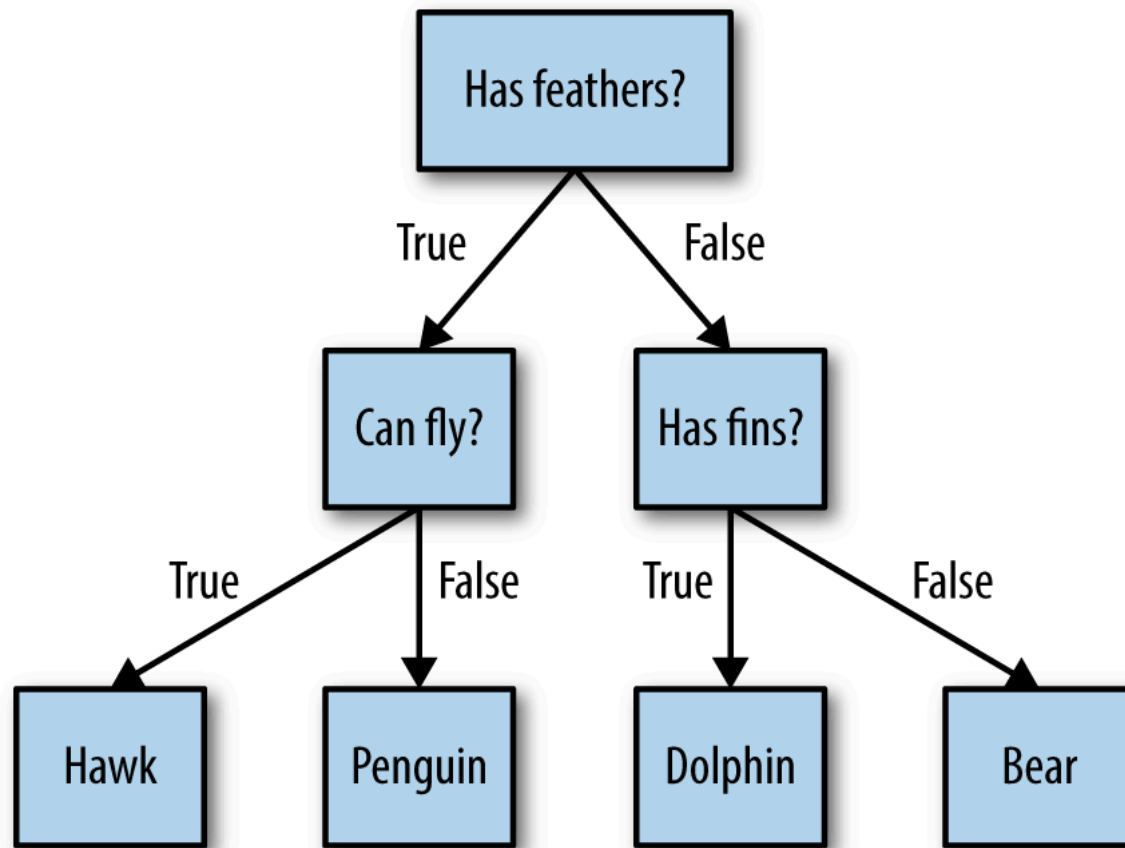


Figure 2-22. A decision tree to distinguish among several animals

**Sometimes referred
to as CART
(Classification and
Regression Trees)**

**Jupyter Notebook
02-supervised-learning
.ipynb [55]**

Decision Trees

- ▶ **Algorithm searches through all possible tests to separate the data into classes**




Figure 2-23. Two-moons dataset on which the decision tree will be built

Decision Trees

- ▶ At each step the algorithm selects the test that provides the best separation of the classes




Figure 2-24. Decision boundary of tree with depth 1 (left) and corresponding tree (right)

Decision Trees

- At each step the algorithm selects the test that provides the best separation of the classes

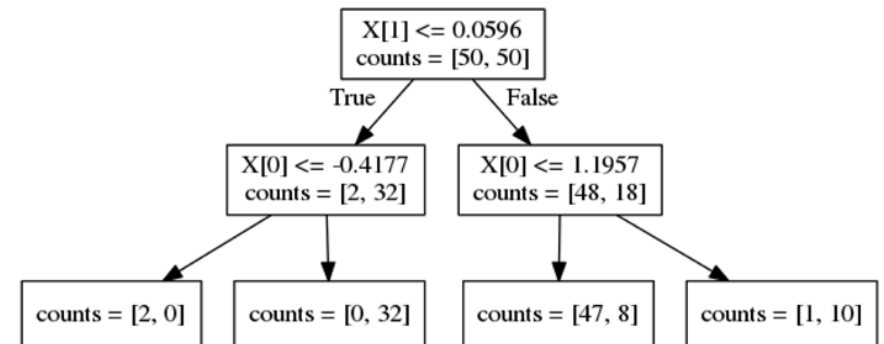



Figure 2-25. Decision boundary of tree with depth 2 (left) and corresponding decision tree (right)

Decision Trees

- ▶ Partitioning is repeated until each “leaf” contains only a single target (class or regression value) - referred to as pure

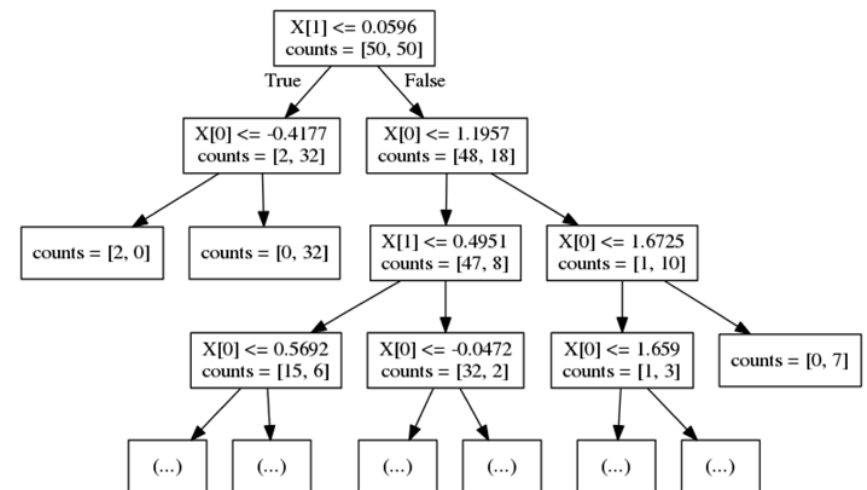



Figure 2-26. Decision boundary of tree with depth 9 (left) and part of the corresponding tree (right); the full tree is quite large and hard to visualize

Decision Trees

- ▶ **Comprehensive tree building leads to overfitting the training data, this can be minimized by**
 - Pre-pruning - stopping the tree building early
 - Can be achieved by limiting the depth of the tree, number of leaves, or only splitting those with a certain number of points
 - Post-pruning/pruning - removing nodes that don't contain much information

**Jupyter Notebook
02-supervised-learning
.ipynb [56-57]**

Decision Trees

- Visualizing the tree can provide information about how the algorithm makes decisions

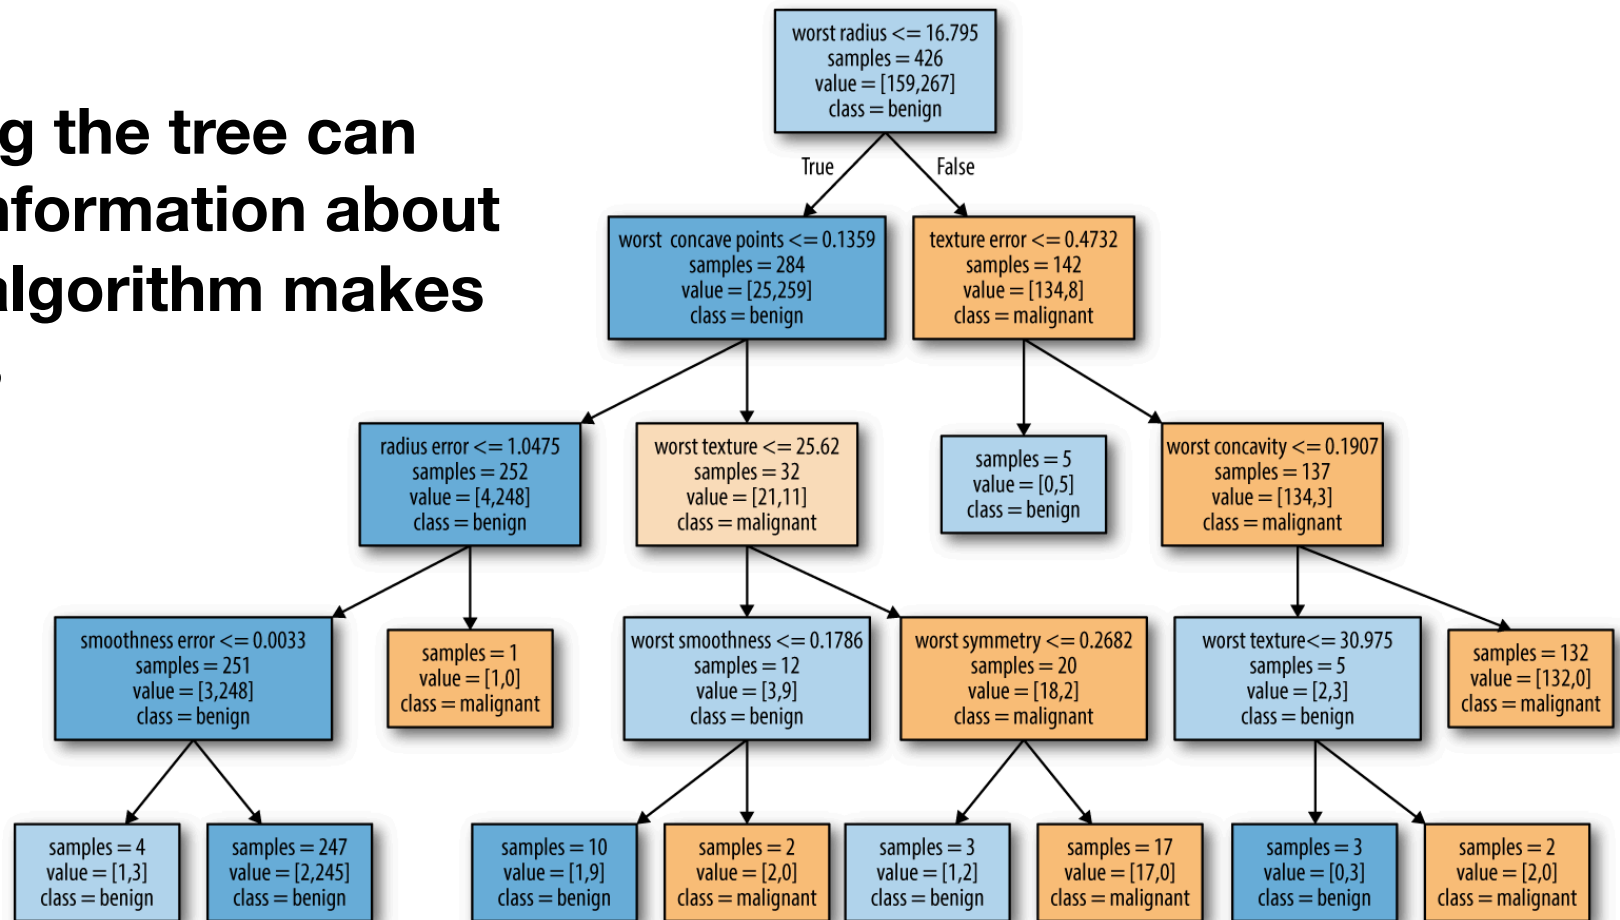


Figure 2-27. Visualization of the decision tree built on the Breast Cancer dataset

Jupyter Notebook
02-supervised-learning
.ipynb [58-59]

Decision Trees

- ▶ Feature importance is a summary of how important each feature is in the tree's decision making
- ▶ Feature importance isn't class specific





Figure 2-28. Feature importances computed from a decision tree learned on the Breast Cancer dataset

Jupyter Notebook
02-supervised-learning
.ipynb [60-61]

Decision Trees Regression



- ▶ Not able to make predictions outside of the range of the training data (extrapolate)

Figure 2-31. Historical development of the price of RAM, plotted on a log scale




Figure 2-32. Comparison of predictions made by a linear model and predictions made by a regression tree on the RAM price data

**Jupyter Notebook
02-supervised-learning
.ipynb [63-65]**

Decision Trees

▸ **Parameters**

- Maximum depth - for pre-pruning

▸ **Assumptions**

- No assumptions about the distribution of the data

▸ **Best Practices**

- Use of pre-pruning prevents overfitting

▸ **Strengths**

- Can be visualized, which aids in interpretation
- Invariant to scaling data
- Works well with mixed data (e.g., binary and continuous features)

▸ **Weaknesses**

- Tend to overfit (even with pre-pruning) and thus don't always generalize
- Doesn't work well on high-dimensional sparse data

Random Forests

- ▶ **Ensemble of slightly different decision trees**
 - Averaging the predictions of the trees minimizing overfitting
 - Difference between trees is achieved by randomizing which data points or features are used

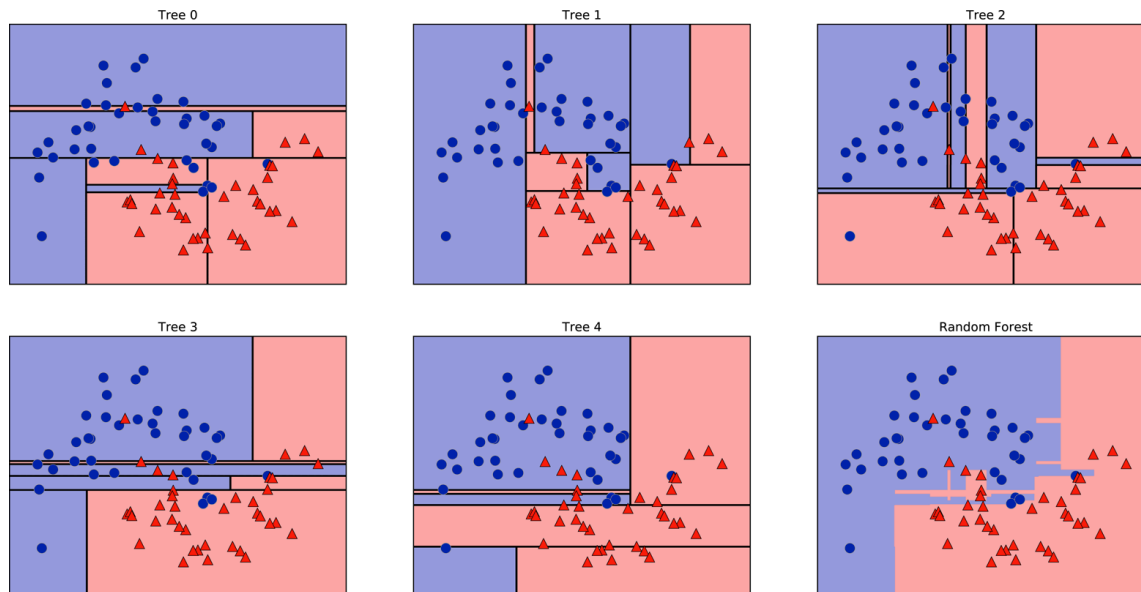


Figure 2-33. Decision boundaries found by five randomized decision trees and the decision boundary obtained by averaging their predicted probabilities

**Jupyter Notebook
02-supervised-learning
.ipynb [66-67]**

Random Forests

- ▶ **Feature importance is aggregated across the ensemble of decision trees**




Figure 2-28. Feature importances computed from a decision tree learned on the Breast Cancer dataset

- Typically more informative for random forests than for individual trees




Figure 2-34. Feature importances computed from a random forest that was fit to the Breast Cancer dataset

Jupyter Notebook
02-supervised-learning
.ipynb [68-69]

Random Forests

▸ Parameters

- Number of estimators - larger is always better
- Number of data points - number of samples drawn from training dataset
- Maximum features - amount of randomness (smaller reduces overfitting)
- Maximum depth - for pre-pruning

▸ Strength

- Invariant to scaling data
- Works well with mixed data (e.g., binary and continuous features)
- Generalizes better than decision trees

▸ Weaknesses

- Harder to interpret than decision trees
- Can be time consuming to train
- Random process can make reproducibility difficult

Gradient Boosted Random Forests

- ▶ Series of decision trees where each new tree tries to correct the mistakes of the previous trees using pre-pruning (rather than randomness)
- ▶ Trees are typically shallow (weak learners)
- ▶ Feature importances tend to be sparser than random forests




Figure 2-35. Feature importances computed from a gradient boosting classifier that was fit to the Breast Cancer dataset

Jupyter Notebook
02-supervised-learning
.ipynb [70-73]

Gradient Boosted Random Forests

▸ **Best Practices**

- Random forests are a good place to start
 - Gradient-boosting can improve accuracy

▸ **Parameters**

- Number of estimators - larger is always better
- Maximum depth - for pre-pruning
- Learning rate - how much a tree tries to correct the previous mistakes

▸ **Strengths**

- Typically perform very well
- Invariant to scaling data
- Works well with mixed data (e.g., binary and continuous features)

▸ **Weaknesses**

- Can take a long time to train
- Doesn't work well on high-dimensional sparse data